

CSCI 544

Applied Natural Language Processing

Mohammad Rostami
USC Computer Science Department

Primarily based on slides prepared by I-Hung Hsu



Logistical Notes

- **Project:**
 - 259 students: 52 groups with one group with 4 members
 - 14 unassigned students
 - Proposal deadline: 3 weeks
- **Written Assignment:**
 - Start peer grading soon after the deadline

PYTORCH

- It's a python-based scientific computing package
- A library that helps using the power of GPUs
- Why popular?
- PyTorch vs TensorFlow

Installation

- Follow instructions in <https://pytorch.org/get-started/locally/>

PyTorch Build	Stable (1.9.0)	Preview (Nightly)	LTS (1.8.2)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python	C++ / Java		
Compute Platform	CUDA 10.2	CUDA 11.1	ROCm 4.2 (beta)	CPU
Run this Command:	conda install pytorch torchvision torchaudio cudatoolkit=10.2 -c pytorch			

PyTorch

PyTorch data structure format

- Tensor: similar to Numpy array but runs on GPU or other hardware accelerators.
- (Trainable) Tensor: stores data and gradient and can be considered similar to a variable

$$y = \beta^T x$$

- Module: A neural network layer; may store state or learnable weights

Tensors vs Numpy Arrays

PyTorch Tensors are just like numpy arrays, but they can run on GPU.

Can have high dimensions

Tensor operations: `x.mm()`, `torch.add(x,y)`, etc

Numpy Conversion: `x.numpy()`

Example

Here a two-layer net is trained using PyTorch Tensors.


- Forward Pass

$$\mathcal{L} = \|y - \hat{y}\|_2^2$$

$$y = W_2^T \text{ReLU}(W_1^T x)$$

- Backword Pass

$$W^{i+1} = W^i - \eta \frac{d\mathcal{L}}{dW^i}$$
$$\frac{d\mathcal{L}}{dW_2} = 2(y - \hat{y}) \frac{d\hat{y}}{dW_2}$$



Optimization with Tensors

Create random tensor for data and weight

Tensor can also be loaded by:

1. Load from data (list)

```
data = [[1, 2], [3, 4]]  
x_data = torch.tensor(data)
```

2. From Numpy array

```
np_array = np.array(data)  
x_np = torch.from_numpy(np_array)
```

3. From another tensor

```
np_array = np.array(data)  
x_np = torch.from_numpy(np_array)
```

```
import torch  
  
dtype = torch.cuda.FloatTensor  
N, D_in, H, D_out = 64, 1000, 100, 10  
x = torch.randn(N, D_in).type(dtype)  
y = torch.randn(N, D_out).type(dtype)  
w1 = torch.randn(D_in, H).type(dtype)  
w2 = torch.randn(H, D_out).type(dtype)
```

```
lr = 1e-6  
for t in range(500):  
    h = x.mm(w1)  
    h_relu = h.clamp(min=0)  
    y_pred = h_relu.mm(w2)  
    loss = (y_pred - y).pow(2).sum()  
  
    grad_y_pred = 2.0 * (y_pred - y)  
    grad_w2 = h_relu.t().mm(grad_y_pred)  
    grad_h_relu = grad_y_pred.mm(w2.t())  
    grad_h = grad_h_relu.clone()  
    grad_h[h < 0] = 0  
    grad_w1 = x.t().mm(grad_h)  
  
    w1 -= lr * grad_w1  
    w2 -= lr * grad_w2
```


Optimization with Tensors

Forward pass: compute predictions and loss



```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

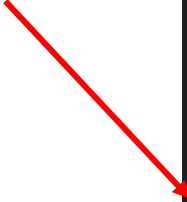
lr = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred-y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred-y)
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= lr*grad_w1
    w2 -= lr*grad_w2
```

Optimization with Tensors

Backward pass: manually computed gradients *if you don't have autograd.*



```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

lr = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred-y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred-y)
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= lr*grad_w1
    w2 -= lr*grad_w2
```

Optimization with Tensors

```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

lr = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred-y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred-y)
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)


    w1 -= lr*grad_w1
    w2 -= lr*grad_w2
```

Optimization: Gradient descent step on weights



Optimization with Tensors

To run on GPU, just cast tensors to a cuda datatype



```
import torch
dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

lr = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()

    grad_y_pred = 2.0 * (y_pred - y)
    grad_w2 = h_relu.t().mm(grad_y_pred)
    grad_h_relu = grad_y_pred.mm(w2.t())
    grad_h = grad_h_relu.clone()
    grad_h[h < 0] = 0
    grad_w1 = x.t().mm(grad_h)

    w1 -= lr * grad_w1
    w2 -= lr * grad_w2
```

Autograd

The previous process:

- Slow
- Gradient is hard to compute when model becomes more complex

=> **This is a primary reason for popularity of PyTorch**

Trainable Tensors:

Set “required_grad” to True to enable torch.autograd

Optimization with Autograd

We set the weights'
“requires_grad” to be True


```
dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)

w3 = w1.detach().clone()
w3.requires_grad=True
# To initialize, you can use:
# w3 = torch.randn(D_in, H, requires_grad=True).type(dtype)

w4 = w2.detach().clone()
w4.requires_grad=True
```

Optimization with Autograd

Forward pass looks exactly the same as the Tensor/Numpy version.



```
lr = 1e-6
for t in range(500):
    y_pred2 = x.mm(w3).clamp(min=0).mm(w4)
    loss = (y_pred2-y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w3 -= lr * w3.grad
        w4 -= lr * w4.grad

    w3.grad.zero_()
    w4.grad.zero_()
```

Optimization with Autograd

But the gradient of loss with respect to w_3 and w_4 can be done by a simple one-line code.

```
lr = 1e-6
for t in range(500):
    y_pred2 = x.mm(w3).clamp(min=0).mm(w4)
    loss = (y_pred2-y).pow(2).sum()
    loss.backward()

    with torch.no_grad():
        w3 -= lr * w3.grad
        w4 -= lr * w4.grad

    w3.grad.zero_()
    w4.grad.zero_()
```

```
grad_y_pred = 2.0 * (y_pred-y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)
```


Optimization with Autograd

Make gradient step on weights.

What's `torch.no_grad()`?

- We need to use `NO_GRAD` to keep the update out of the gradient computation
- Why is that? It boils down to the DYNAMIC GRAPH that PyTorch uses.

```
lr = 1e-6
for t in range(500):
    y_pred2 = x.mm(w3).clamp(min=0).mm(w4)
    loss = (y_pred2 - y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w3 -= lr * w3.grad
        w4 -= lr * w4.grad

    w3.grad.zero_()
    w4.grad.zero_()
```

Optimization with Autograd

grad.zero()?

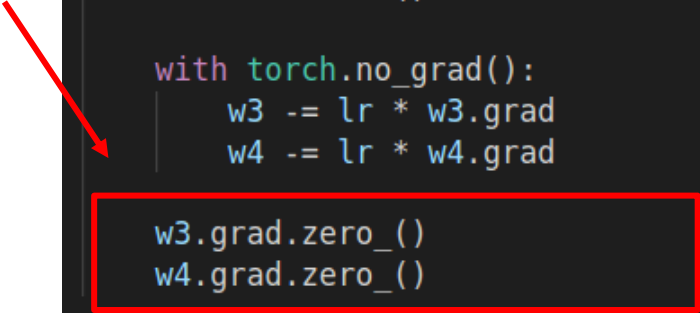
- After each gradient descent step, we should restart from zero
- Why is that? Gradients are accumulated

```
lr = 1e-6
for t in range(500):
    y_pred2 = x.mm(w3).clamp(min=0).mm(w4)
    loss = (y_pred2-y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w3 -= lr * w3.grad
        w4 -= lr * w4.grad

    w3.grad.zero_()
    w4.grad.zero_()
```



New Autograd Functions

We can define our own autograd functions by writing forward and backward for Tensors

```
class ReLU(torch.autograd.Function):  
  
    @staticmethod  
    def forward(ctx, x):  
        ctx.save_for_backward(x)  
        return x.clamp(min=0)  
  
    @staticmethod  
    def backward(ctx, grad_y):  
        x, = ctx.saved_tensors  
        grad_input = grad_y.clone()  
        grad_input[x < 0] = 0  
        return grad_input
```

New Autograd Functions

```
class ReLU(torch.autograd.Function):
```

```
    @staticmethod
```

```
    def forward(ctx, x):
```

```
        ctx.save_for_backward(x)
```

```
        return x.clamp(min=0)
```

```
    @staticmethod
```

```
    def backward(ctx, grad_y):
```

```
        x, = ctx.saved_tensors
```

```
        grad_input = grad_y.clone()
```

```
        grad_input[x < 0] = 0
```

```
        return grad_input
```

```
lr = 1e-6
```

```
for t in range(500):
```

```
    y_pred3 = ReLU.apply(x.mm(w5)).mm(w6)
```

```
    loss = (y_pred3 - y).pow(2).sum()
```

```
    loss.backward()
```

```
    with torch.no_grad():
```

```
        w5 -= lr * w5.grad
```

```
        w6 -= lr * w6.grad
```

```
    w5.grad.zero_()
```

```
    w6.grad.zero_()
```

We then apply our new function in the forward pass.

Networks: torch.nn

- High-level wrapper for creating neural networks
- This is another reason behind popularity of PyTorch and TensorFlow
- Various classes of neural networks can be built using built-in modules, e.g., CNN, RNN, LSTM, transformers, etc
- A diverse set of hyperparameters are built-in implemented, e.g., various activation function, loss functions, etc.

torch.nn

Define our model as a
sequence of layers

```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out)
)
model.cuda()

loss_fn = torch.nn.MSELoss(reduction='sum')

def weights_init(m):
    if isinstance(m, torch.nn.Linear):
        torch.nn.init.zeros_(m.weight)
        torch.nn.init.ones_(m.bias)

model.apply(weights_init)
print(model)

lr = 1e-6

for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    model.zero_grad()
    loss.backward()

    with torch.no_grad():
        for param in model.parameters():
            param.data -= lr * param.grad.data
```

NN also can be used for
common loss functions

torch.nn

Forward pass:

- Feed data to model
- Use prediction to and ground truth to get loss function

```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out)
)
model.cuda()

loss_fn = torch.nn.MSELoss(reduction='sum')

def weights_init(m):
    if isinstance(m, torch.nn.Linear):
        torch.nn.init.zeros_(m.weight)
        torch.nn.init.ones_(m.bias)

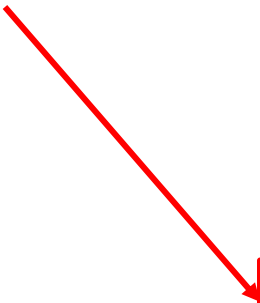
model.apply(weights_init)
print(model)

lr = 1e-6

for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    model.zero_grad()
    loss.backward()

    with torch.no_grad():
        for param in model.parameters():
            param.data -= lr * param.grad.data
```



torch.nn

Pytorch generates autograd easily

We should use `model.zero_grad()`
to clear all gradients for the
parameters in the model

```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out)
)
model.cuda()

loss_fn = torch.nn.MSELoss(reduction='sum')

def weights_init(m):
    if isinstance(m, torch.nn.Linear):
        torch.nn.init.zeros_(m.weight)
        torch.nn.init.ones_(m.bias)

model.apply(weights_init)
print(model)

lr = 1e-6

for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    model.zero_grad()
    loss.backward()

    with torch.no_grad():
        for param in model.parameters():
            param.data -= lr * param.grad.data
```


Optimizer

Make gradient step on each model parameter.

Question:

- How can we apply more advanced rules for updating

```
import torch

dtype = torch.cuda.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out)
)
model.cuda()

loss_fn = torch.nn.MSELoss(reduction='sum')

def weights_init(m):
    if isinstance(m, torch.nn.Linear):
        torch.nn.init.zeros_(m.weight)
        torch.nn.init.ones_(m.bias)

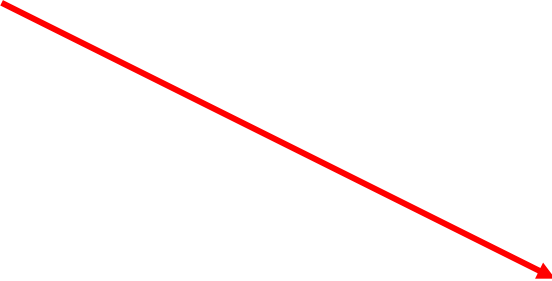
model.apply(weights_init)
print(model)

lr = 1e-6

for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    model.zero_grad()
    loss.backward()


    with torch.no_grad():
        for param in model.parameters():
            param.data -= lr * param.grad.data
```



Optimizer

Call nn.optim package, which contains various advanced optimizer other than SGD.

Now, all the parameters can be updated via one-line code.



```
optimizer = torch.optim.Adam(model.parameters(), lr=lr)
for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    #model.zero_grad()
    optimizer.zero_grad()
    loss.backward()

    optimizer.step()
# with torch.no_grad():
#     for param in model.parameters():
#         param.data -= lr * param.grad.data
```

Define new modules

Pytorch **Module** is
a neural network
layer, it can contain
weights or other
modules.

```
class TwoLayerMLP(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super().__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)
        self.weight_init()

    def weight_init(self):
        torch.nn.init.zeros_(self.linear1.weight)
        torch.nn.init.zeros_(self.linear2.weight)
        torch.nn.init.ones_(self.linear1.bias)
        torch.nn.init.ones_(self.linear2.bias)

    def forward(self, x):
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred

model = TwoLayerMLP(D_in, H, D_out)
model.cuda()
print(model)
loss_fn = torch.nn.MSELoss(reduction='sum')
lr = 1e-6

optimizer = torch.optim.Adam(model.parameters(), lr=lr)
for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

DataLoaders

A **DataLoader** wraps a **Dataset** and provides minibatching, multithreading, etc.

When you need to load custom data, just **write your own Dataset class**

```
import torch
from torch.utils import data

class Dataset(data.Dataset):
    'Characterizes a dataset for PyTorch'
    def __init__(self, list_IDs, labels):
        'Initialization'
        self.labels = labels
        self.list_IDs = list_IDs

    def __len__(self):
        'Denotes the total number of samples'
        return len(self.list_IDs)


    def __getitem__(self, index):
        'Generates one sample of data'
        # Select sample
        ID = self.list_IDs[index]

        # Load data and get label
        X = torch.load('data/' + ID + '.pt')
        y = self.labels[ID]


        return X, y
```

(Mini)batching

All your
training data:
Standard for
loop:



(Mini)
batching:



Adapt Dataset to DataLoaders

```
# Parameters
params = {'batch_size': 64,
          'shuffle': True,
          'num_workers': 6}
max_epochs = 100

# Datasets
partition = # IDs
labels = # Labels

# Generators
training_set = Dataset(partition['train'], labels)
training_generator = data.DataLoader(training_set, **params)

validation_set = Dataset(partition['validation'], labels)
validation_generator = data.DataLoader(validation_set, **params)

# Loop over epochs
for epoch in range(max_epochs):
    # Training
    for local_batch, local_labels in training_generator:
        # Transfer to GPU
        local_batch, local_labels = local_batch.to(device), local_labels.to(device)

        # Model computations
        [...]
```

**DataLoader perform
batching
automatically**

Summary

1. Prepare you data
 - a. Write your own Dataset (inherit `torch.nn.util.dataset`)
2. Create your model
 - a. A sequential module if your model is super easy and will not be reused.
 - b. A `nn.Module` module
3. Write the loop (how many epoch/steps) to train your model:
 - a. Create a `DataLoader` that wraps the Dataset you provide
 - b. Set an optimizer
 - c. Set a loss for optimization
 - d. For loop....
 - i. Forward pass
 - ii. Zero-grad
 - iii. Backward pass => Get gradient
 - iv. Optimizer step to update your models' weight.

Using Different Version of Data for Efficiency

- Tiny-size: for debugging syntactic bug
- Small-size: check the behavior of the model
- Mid-size: for understanding model behavior and fast development
- Full-size: conduct final experiments

Tensorboard

- Installation:

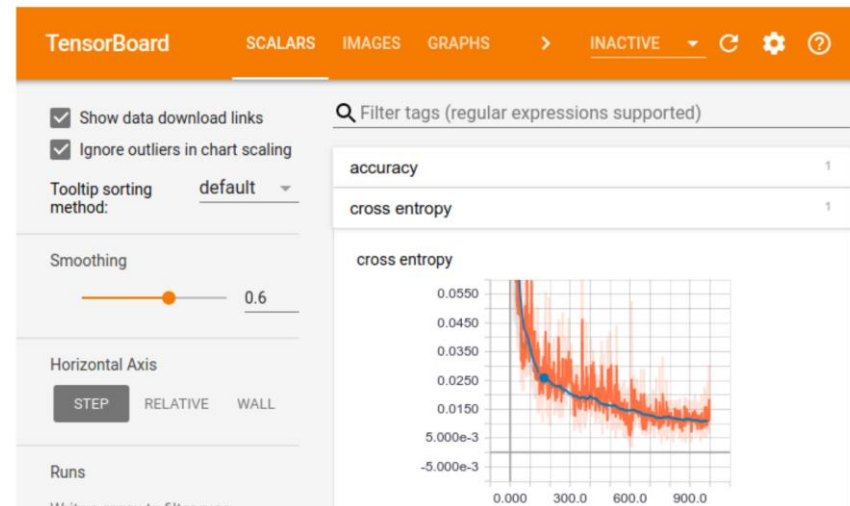
https://pytorch.org/tutorials/recipes/recipes/tensorboard_with_pytorch.html

- Create a Summary Writer

```
from torch.utils.tensorboard import SummaryWriter  
writer = SummaryWriter()
```

- During Training/Dev/Test

```
writer.add_scalar("Loss/train", loss, epoch)
```



End to end Examples

- <https://www.analyticsvidhya.com/blog/2020/01/first-text-classification-in-pytorch/>
- <https://towardsdatascience.com/lstm-text-classification-using-pytorch-2c6c657f8fc0>